

Business Data Mining

IDS 472 (Spring 2024)


Instructor: Wenxin Zhou

- In this chapter, we formalize the kNN mechanism for both **classification** and **regression**

NONPARAMETRIC DISCRIMINATION: CONSISTENCY PROPERTIES

1. Introduction

The discrimination problem (two population case) may be defined as follows: a random variable Z , of observed value z , is distributed over some space (say, p -dimensional) either according to distribution F , or according to distribution G . The problem is to decide, on the basis of z , which of the two distributions Z has.



- Binary classification: $y \in \{0, 1\}$
- Training set: $\mathbf{S}_{tr} = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n)\}$
- Given a feature vector \mathbf{x} (test observation), denote the i^{th} **nearest** observation to \mathbf{x} as $\mathbf{x}_{(i)}(\mathbf{x})$ with label $y_{(i)}(\mathbf{x})$
- Standard kNN classifier $\psi(\mathbf{x})$ is given by

$$\hat{y} = \psi(\mathbf{x}) = \begin{cases} 1 & \text{if } \sum_{i=1}^k \frac{1}{k} I_{\{y_{(i)}(\mathbf{x})=1\}} > \sum_{i=1}^k \frac{1}{k} I_{\{y_{(i)}(\mathbf{x})=0\}}, \\ 0 & \text{otherwise,} \end{cases}$$

where k is the number of nearest neighbors wrt \mathbf{x}

- For multiclass classification with c classes, $y \in \{0, 1, \dots, c - 1\}$, the kNN classifier is

$$\hat{y} = \psi(\mathbf{x}) = \operatorname{argmax}_j \sum_{i=1}^k \frac{1}{k} I_{\{y_{(i)}(\mathbf{x})=j\}}$$

- Avoid an even k to **prevent ties** (3NN, 5NN)
- The **larger k** , the **smoother** are the *decision boundaries* (boundaries between *decision regions*)
- Below we examine the effect of k on the decision boundaries and accuracy of kNN on scaled Iris data

```
arrays = np.load('data/iris_train_scaled.npz')
X_train = arrays['X']
y_train = arrays['y']
arrays = np.load('data/iris_test_scaled.npz')
X_test = arrays['X']
y_test = arrays['y']
```

```
print('X shape = {}'.format(X_train.shape) \
      + '\ny shape = {}'.format(y_train.shape))
```

✓ 0.0s

Python

X shape = (120, 4)

y shape = (120,)

- For illustration purposes, consider the first two features in data

```
X_train = X_train[:,[0,1]]  
X_test = X_test[:,[0,1]]  
X_train.shape
```

(120, 2)

- Implement 4 kNN classifiers for $k = 1, 3, 9, 36$
- Visualize the resulting decision regions:
 - great a grid
 - train a classifier using training data
 - classify each point in the grid using trained classifier
 - plot the decision regions based on the assigned labels to each point in the grid

Classification



import required functionality for this chapter

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib.colors import ListedColormap
from sklearn.neighbors import KNeighborsClassifier as KNN
```

```
color = ['aquamarine', 'bisque', 'lightgrey']
cmap = ListedColormap(color)

mins = X_train.min(axis=0) - 0.1
maxs = X_train.max(axis=0) + 0.1
x = np.arange(mins[0], maxs[0], 0.01)
y = np.arange(mins[1], maxs[1], 0.01)
X, Y = np.meshgrid(x, y)
coordinates = np.array([X.ravel(), Y.ravel()]).T

fig, axs = plt.subplots(2, 2, figsize=(6, 4), dpi = 200)
fig.tight_layout()
K_val = [1, 3, 9, 36]
```

```
for ax, K in zip(axes.ravel(), K_val):
    knn = KNN(n_neighbors=K, weights='uniform', metric='euclidean')
    knn.fit(X_train, y_train)
    Z = knn.predict(coordinates)
    Z = Z.reshape(X.shape)
    ax.tick_params(axis='both', labelsize=6)
    ax.set_title(str(K) + 'NN Decision Regions', fontsize=8)
    ax.pcolormesh(X, Y, Z, cmap = cmap, shading='nearest')
    ax.contour(X, Y, Z, colors='black', linewidths=0.5)
    ax.plot(X_train[y_train==0, 0], X_train[y_train==0, 1], 'g.', markersize=4)
    ax.plot(X_train[y_train==1, 0], X_train[y_train==1, 1], 'r.', markersize=4)
    ax.plot(X_train[y_train==2, 0], X_train[y_train==2, 1], 'k.', markersize=4)
    ax.set_xlabel('sepal length (normalized)', fontsize=7)
    ax.set_ylabel('sepal width (normalized)', fontsize=7)

    print('The accuracy for K={} on the training data is {:.3f}'\
          .format(K, knn.score(X_train, y_train)))
    print('The accuracy for K={} on the test data is {:.3f}'\
          .format(K, knn.score(X_test, y_test)))

for ax in axes.ravel():
    # show the x-label and the y-label for the last row and the left column, respectively
    ax.label_outer()
```

similarly to interpolator: low bias, high variance



The accuracy for $K=1$ on the training data is 0.933

The accuracy for $K=1$ on the test data is 0.667

The accuracy for $K=3$ on the training data is 0.833

The accuracy for $K=3$ on the test data is 0.767

The accuracy for $K=9$ on the training data is 0.858

The accuracy for $K=9$ on the test data is 0.800

The accuracy for $K=36$ on the training data is 0.775


The accuracy for $K=36$ on the test data is 0.800




similarly to a smooth curve: larger bias, low variance

Classification

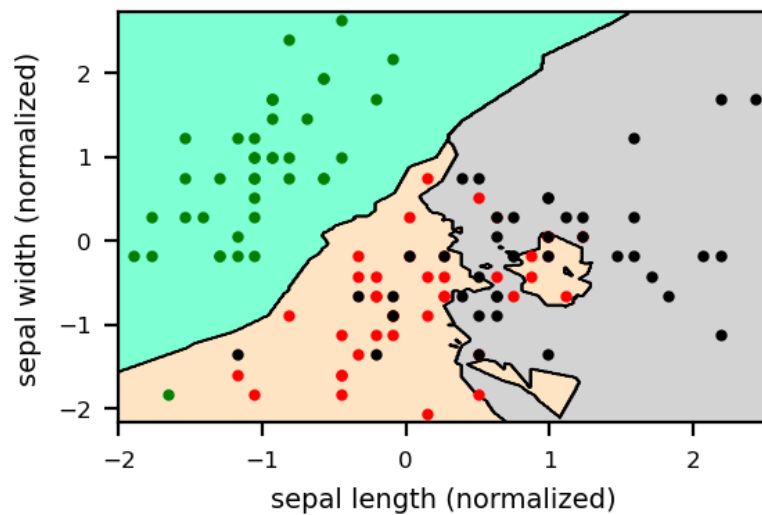
1NN Decision Regions




3NN Decision Regions



9NN Decision Regions



36NN Decision Regions



- Another form of kNN is **distance-weighted kNN (DW-kNN)**
 - the k nearest neighbors of a test observation \mathbf{x} are weighted according to their **distances** from \mathbf{x}
 - observations that are **closer** to \mathbf{x} should impose a **higher influence** on decision making
 - use the **inverse** of distance to \mathbf{x} as weight

$$\psi(\mathbf{x}) = \begin{cases} 1 & \text{if } \frac{1}{D} \sum_{i=1}^k \frac{1}{d[\mathbf{x}_{(i)}(\mathbf{x}), \mathbf{x}]} I_{\{y_{(i)}(\mathbf{x})=1\}} > \frac{1}{D} \sum_{i=1}^k \frac{1}{d[\mathbf{x}_{(i)}(\mathbf{x}), \mathbf{x}]} I_{\{y_{(i)}(\mathbf{x})=0\}}, \\ 0 & \text{otherwise,} \end{cases}$$

$D = \sum_{i=1}^k \frac{1}{d[\mathbf{x}_{(i)}(\mathbf{x}), \mathbf{x}]}$

where $d[\mathbf{x}_{(i)}(\mathbf{x}), \mathbf{x}]$ is the distance of $\mathbf{x}_{(i)}(\mathbf{x})$ from \mathbf{x}

- In practice, the choice of using DW-kNN, standard kNN, or the choice of k will be decided in the **model selection phase**

- DW-kNN is implemented by setting `weights='distance'` in `KNeighborsClassifier` (default is `'uniform'`)





Fig. 1. Plots of probability of error with respect to k for the two nearest-neighbor type rules.

Choice of Distance

- Let $\mathbf{x}_i = [x_{1i}, x_{2i}, \dots, x_{qi}]^T$ and $\mathbf{x}_j = [x_{1j}, x_{2j}, \dots, x_{qj}]^T$
- Three popular choices of distances are Euclidean, Manhattan and Minkowski



$$d_E[\mathbf{x}_i, \mathbf{x}_j] = \sqrt{\sum_{l=1}^q (x_{li} - x_{lj})^2}$$

$$d_{Ma}[\mathbf{x}_i, \mathbf{x}_j] = \sum_{l=1}^q |x_{li} - x_{lj}|$$

$$d_{Mi}[\mathbf{x}_i, \mathbf{x}_j] = \left(\sum_{l=1}^q |x_{li} - x_{lj}|^p \right)^{\frac{1}{p}}$$

- In `KNeighborsClassifier`, the choice of distance is determined by the `metric` parameter

Manhattan vs Euclidean distance



- The kNN regressor $f(\mathbf{x})$ is given by

$$\hat{y} = f(\mathbf{x}) = \sum_{i=1}^k \frac{1}{k} y_{(i)}(\mathbf{x})$$

- The standard kNN regressor estimates the target of a given \mathbf{x} as the **average of k targets** of the nearest neighbors of \mathbf{x}
- Hereafter, denote regressors simply as $f(\mathbf{x})$
- In this course, a regressor refers to an estimator of the *conditional mean function*

$$E(Y|X = \mathbf{x}), \mathbf{x} \in \mathbb{R}^p$$

A Regression Application

- California Housing dataset
 - median house price (in \$100,000) of 20640 CA districts
 - use 8 features to predict the house price
 - the target is the median house price of a district ('MedHouseVal' in the data)

```
from sklearn import datasets
from sklearn.neighbors import KNeighborsRegressor as KNN

california = datasets.fetch_california_housing()
print('california housing data shape: ' + str(california.data.shape) + \
      '\nfeature names: ' + str(california.feature_names) + \
      '\ntarget name: ' + str(california.target_names))
```

✓ 0.0s

Python

```
california housing data shape: (20640, 8)
feature names: ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms', 'Population', 'AveOccup', 'Latitude', 'Longitude']
target name: ['MedHouseVal']
```

A Regression Application

- Use the DESCR key to check some details about the dataset

```
print(california.DESCR[:975])
```

Python

```
.. _california_housing_dataset:
```

```
California Housing dataset
```

```
-----
```

```
**Data Set Characteristics:**
```

```
:Number of Instances: 20640
```

```
:Number of Attributes: 8 numeric, predictive attributes and the target
```

```
:Attribute Information:
```

```
- MedInc      median income in block group
- HouseAge    median house age in block group
- AveRooms    average number of rooms per household
- AveBedrms   average number of bedrooms per household
- Population  block group population
- AveOccup    average number of household members
- Latitude    block group latitude
- Longitude   block group longitude
```

```
:Missing Attribute Values: None
```

This dataset was obtained from the StatLib repository.

https://www.dcc.fc.up.pt/~ltorgo/Regression/cal_housing.html

The target variable is the median house value for California districts, expressed in hundreds of thousands of dollars (\$100,000).

- Split the data into **training** and **test** sets: use the default 0.25 `test_size` and do not set `stratify` to any variable (there are no “classes” in regression)

```
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test= train_test_split(california.data,
                                                    california.target,
                                                    random_state=100)
print('X_train_shape: ' + str(X_train.shape) + '\nX_test_shape: ' \
      + str(X_test.shape) + '\ny_train_shape: ' + str(y_train.shape)\
      + '\ny_test_shape: ' + str(y_test.shape))
```

Python

```
X_train_shape: (15480, 8)
X_test_shape: (5160, 8)
y_train_shape: (15480,)
y_test_shape: (5160,)
```

- Use standardization to scale the **training** and **test** sets: the scaler object is trained using the training set, and is then used to transform both training and test sets

```
scaler = StandardScaler()  
scaler.fit(X_train)  
X_train_scaled = scaler.transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

- Previously we used **pair plots** as an exploratory analysis.
- Another exploratory analysis is to look into **Pearson's correlation coefficients** between predictors and the response or even between predictors themselves
- The Pearson's correlation coefficient between two rv X and Y is

$$\rho = \frac{\text{Cov}[X, Y]}{\sigma_X \sigma_Y}$$

$$\text{Cov}[X, Y] = E[(X - E[X])(Y - E[Y])]$$

- Given a sample $\{(x_1, y_1), \dots, (x_n, y_n)\}$ from (X, Y) . The sample Pearson's correlation coefficient is

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

- sample correlation coefficient is **scale and location invariant**
 - it is a measure of the degree of **linear relationship**
 - $-1 \leq \rho \leq 1, -1 \leq r \leq 1$
 - $|\rho| < 1$ indicates the relationship is not completely linear
- `pandas.DataFrame.corr()` calculates pairwise correlation coefficients

Exploratory Analysis

```
# np.concatenate((a1, a2)) concatenate arrays
# a1 and a2 along the specified axis
california_arr = np.concatenate((X_train_scaled, y_train.reshape(-1,1)),
                                axis=1)
california_pd = pd.DataFrame(california_arr,
                              columns=[*california.feature_names, 'MEDV'])
california_pd.head().round(3)
```


Python

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude	MEDV
0	1.604	-0.843	0.974	-0.086	0.068	0.027	-0.865	0.883	2.903
1	-0.921	0.345	-0.197	-0.238	-0.472	-0.068	1.647	-0.999	0.687
2	-0.809	1.849	-0.376	-0.037	-0.516	-0.082	1.675	-0.741	1.097
3	0.597	-0.289	-0.437	-0.119	-0.680	-0.121	1.008	-1.423	4.600
4	0.219	0.107	0.187	0.122	-0.436	-0.057	0.956	-1.283	2.134

```
fig, ax = plt.subplots(figsize=(9,5))
sns.heatmap(california_pd.corr().round(2), annot=True, square=True, ax=ax)
```

Exploratory Analysis

- Calculate and draw a color-coded plot (**heatmap**) of these correlation coefficients using `seaborn.heatmap()`



- One way that we may use this exploratory analysis is in *feature selection*
- In many applications with a **limited sample size** and **moderate to large number of features**, using a **subset** of features could lead to a better performance in predicting the target than using all features
- This is due to what is known as the “**curse of dimensionality**”
- Using more features in training not only increases the computational burden, but also could lead to a lower performance of the trained models after adding more than a certain number of features
- We may use the **correlation matrix** to identify a subset of features such that each feature within this subset is **strongly** or **moderately** correlated with the response

- Defining a strong or moderate correlation is subjective
- For simplicity, we choose variables with a correlation magnitude ≥ 0.1 : MedInc, HouseAge, AveRooms, Latitude

```
# X_trained with selected features
X_train_fs_scaled = X_train_scaled[:, [0, 1, 2, 7]]
X_test_fs_scaled = X_test_scaled[:, [0, 1, 2, 7]]
print('The shape of training X after feature selection: '\
      + str(X_train_fs_scaled.shape))
print('The shape of test X after feature selection: '\
      + str(X_test_fs_scaled.shape))
```

✓ 0.0s

Python

The shape of training X after feature selection: (15480, 4)

The shape of test X after feature selection: (5160, 4)

```
# save for possible uses later
np.savez('data/california_train_fs_scaled',
         X = X_train_fs_scaled, y = y_train)
np.savez('data/california_test_fs_scaled',
         X = X_test_fs_scaled, y = y_test)
```

✓ 0.0s


Python

- The standard kNN regressor is implemented in the `KNeighborsRegressor` class from `sklearn.neighbors` module
- Train a standard 50NN (Euclidean distance & uniform weights) using the scaled training set (`X_train_fs_scaled`), then use the model to predict the target in the test set (`X_test_fs_scaled`)
- Create scatter plot between predicted targets \hat{y} and actual targets

```
plt.style.use('seaborn-v0_8')
knn = KNN(n_neighbors=50, weights='uniform', metric='euclidean')
knn.fit(X_train_fs_scaled, y_train)
y_test_predictions = knn.predict(X_test_fs_scaled)

plt.figure(figsize=(4.5, 3), dpi = 200)
plt.plot(y_test, y_test_predictions, 'g.', markersize=4)
lim_left, lim_right = plt.xlim()
plt.plot([lim_left, lim_right], [lim_left, lim_right], '--k', linewidth=1)
plt.xlabel("y (Actual MEDV)", fontsize='small')
plt.ylabel("$\hat{y}$ (Predicted MEDV)", fontsize='small')
plt.tick_params(axis='both', labelsize=7)
```


- The scatter plot of predicted targets by a 50NN and the actual targets in the California Housing dataset



- Let S_{te} be the test set containing m test observations
- If $\hat{y}_i = y_i$ for $\mathbf{x}_i \in S_{te}$, we commit no error in predictions, meaning all points in the scatter plot between predicted and actual targets should lie on the diagonal line $\hat{y} = y$
- From the plot we see that
 - when $MEDV \leq \$100k$, the model generally **overestimates** the target
 - for large values of $MEDV$ around $\$100k$, the model **underestimates** the target


- Classifiers in scikit-learn have a `score` method that given a test data and their labels, returns a **performance measure**
- For regressors, the `score` method estimates **coefficient of determination** (*R-squared statistic* \hat{R}^2), given by

$$\hat{R}^2 = 1 - \frac{\text{RSS}}{\text{TSS}}$$

where RSS and TSS are short for *Residual Sum of Squares* and *Total Sum of Squares*, resp:

$$\text{RSS} = \sum_{i=1}^m (y_i - \hat{y}_i)^2$$

$$\text{TSS} = \sum_{i=1}^m (y_i - \bar{y})^2$$

$$\bar{y} = \frac{1}{m} \sum_{i=1}^m y_i$$


trivial estimator of the target

- \hat{R}^2 measures how well the trained regressor is performing when compared with the trivial estimator
 - for perfect predictions, $RSS = 0$ and $\hat{R}^2 = 1$
 - when trained regressor performs similarly to the trivial one, $RSS = TSS$ and $\hat{R}^2 = 0$

```
print('The test R^2 is: {:.2f}'\n      .format(knn.score(X_test_fs_scaled, y_test)))
```

✓ 0.0s

Python

The test R^2 is: 0.64

- DW-kNN regressor computes a weighted average of targets for the k nearest neighbors, where the weights are the inverse of the distance of each nearest neighbor \mathbf{x}

$$f(\mathbf{x}) = \frac{1}{D} \sum_{i=1}^K \frac{1}{d[\mathbf{x}_{(i)}(\mathbf{x}), \mathbf{x}]} y_{(i)}$$

- Observations that are closer to a test observation impose a higher influence in the weighted average
- In scikit-learn, DW-kNN regressor is obtained by switching the weights parameter of KNeighborsRegressor from 'uniform' to 'distance'